Ensembl gene annotation project (*e!*61)

Mus musculus (mouse, NCBIM37 assembly)

Raw Computes Stage: Searching for sequence patterns, aligning proteins and cDNAs to the genome.

Approximate time: 1 week

The annotation process of the high-coverage mouse assembly began with the raw compute stage [Figure 1] whereby the genomic sequence was screened for sequence patterns including repeats using RepeatMasker [1.] (version 3.2.8 with parameters '-nolow -species "mouse" -s'), Dust [2.] and TRF [3.]. RepeatMasker and Dust combined masked 43.40% of the species genome.
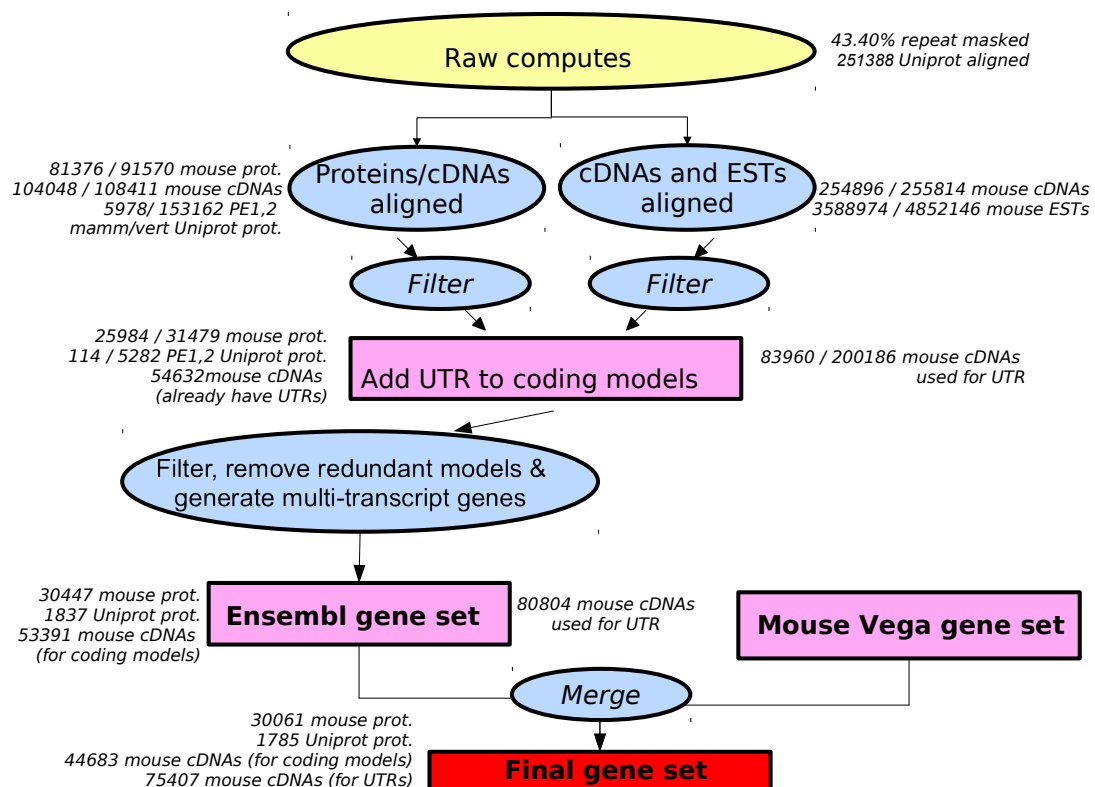

43.40% repeat masked
251388 Uniprot aligned

Raw computes

81376 / 91570 mouse prot.
104048 / 108411 mouse cDNAs
5978/ 153162 PE1,2
mamm/vert Uniprot prot.

Proteins/cDNAs aligned

cDNAs and ESTs aligned

254896 / 255814 mouse cDNAs
3588974 / 4852146 mouse ESTs

Filter

Filter

25984 / 31479 mouse prot.
114 / 5282 PE1,2 Uniprot prot.
54632mouse cDNAs
(already have UTRs)

Add UTR to coding models

83960 / 200186 mouse cDNAs
used for UTR

Filter, remove redundant models & generate multi-transcript genes

30447 mouse prot.
1837 Uniprot prot.
53391 mouse cDNAs
(for coding models)

Ensembl gene set

80804 mouse cDNAs
used for UTR

Mouse Vega gene set

Merge

30061 mouse prot.
1785 Uniprot prot.
44683 mouse cDNAs (for coding models)
75407 mouse cDNAs (for UTRs)

Final gene set

Figure 1: Summary of mouse gene annotation project.

Transcription start sites were predicted using Eponine–scan [4.] and FirstEF [5.]. CpG islands and tRNAs [6.] were also predicted. Genscan [7.] was run across RepeatMasked sequence and the results were used as input for UniProt [8.], UniGene [9.] and Vertebrate RNA [10.] alignments by WU-BLAST [11.]. (Passing only Genscan results to BLAST is an effective way of reducing the search space and therefore the computational resources required.) This resulted in 251388 UniProt, 337375 UniGene and 327209 Vertebrate RNA sequences aligning to the genome.

Targetted Stage: *Generating coding models from mouse evidence*

Approximate time: 6 weeks

Next, mouse protein and cDNA sequences were downloaded from public databases (UniProt SwissProt/TrEMBL [8.] and RefSeq [9.] for proteins, ENA/Genbank/DDBJ and RefSeq [9.] for cDNAs) and filtered to remove sequences based on predictions. The mouse protein sequences were first mapped to rough locations in the genome using Pmatch to reduce the search space for the subsequent Genewise step, as indicated in [Figure 2]. Models of the coding sequence (CDS) were produced from the proteins using Genewise [13.], which was run with four different sets of parameters to accommodate for cases where some coding models contain non-canonical (non GT/AG) splice sites. In parallel to the Genewise step, mouse cDNAs with known CDS start/end coordinates were aligned to the genome using Exonerate (*cdna2genome* model) [12.] to generate coding models [Figure 2]. Because all cDNAs used in this step had known pairing with proteins (e.g. RefSeq cDNAs with accession prefix "NM_" matching RefSeq proteins with "NP_" prefix), it allowed the comparison of coding models generated by Exonerate for a given cDNA to those generated by Genewise using its counterpart protein. All coding models generated by Genewise and Exonerate were filtered systematically by a series of Perl scripts to remove models with erroneous structures (e.g., interlocking models with long introns on the same strand). In addition, models supported by dubious mouse

2

protein/cDNA evidence (e.g. cDNA fragments with wrongly annotated short open-reading frames) were removed manually on a case-by-case basis. The Apollo software [15.] was used to visualise the results of filtering.

Where one protein sequence had generated more than one candidate coding model at a locus, the BestTargetted module was used to select the coding model that most closely matched the source protein to take through to the next stage of the gene annotation process. The generation of transcript models using species-specific (in this case, mouse) data is referred to as the "Targetted stage". This stage resulted in 95192 coding models (36760 built from 31507 mouse proteins and 58432 built from 54632 mouse cDNAs) which were taken through to the UTR addition stage.
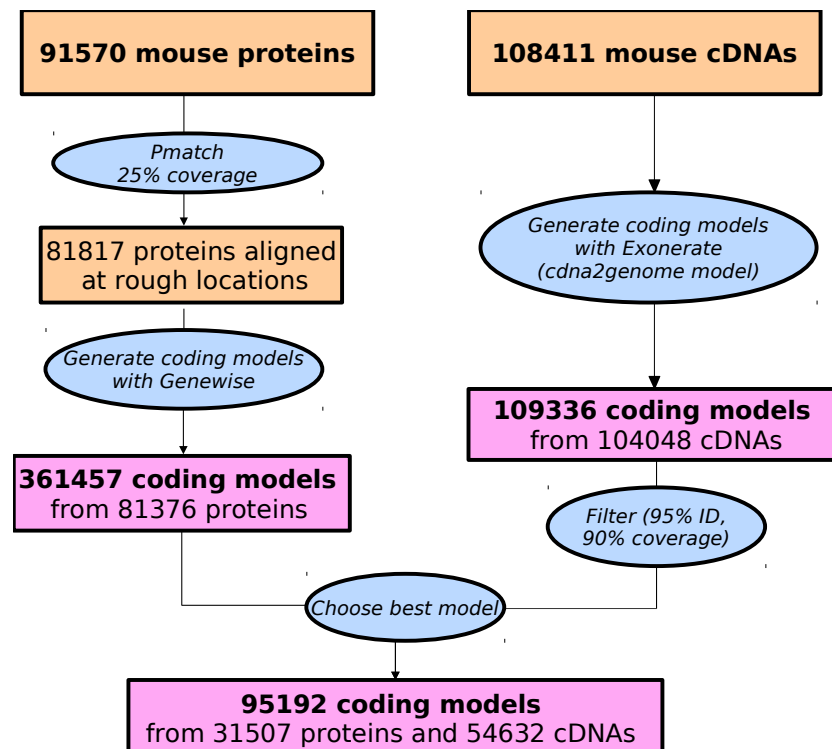

Figure 2: Targetted stage using mouse protein and cDNA sequences.

Similarity Stage: Generating additional coding models using proteins from related species

Approximate time: 2 weeks

Following the mouse Targetted alignments, additional coding models were

generated as follows. The UniProt alignments from the Raw Computes step were filtered to retain only those sequences belonging to UniProt's "Mammalia" and "Vertebrata" taxonomical classes as well as Uniprot's Protein Existence (PE) classification level 1 and 2. In genomic regions which were not covered by any coding models from Targetted alignments, WU-BLAST was rerun for the Uniprot protein sequences and the results were passed to Genewise [13.] to build coding models. In most cases, multiple coding models built from different Uniprot proteins were generated in a single locus, each model with a slightly different exon-intron structure. To filter for the best supported structures, the TranscriptConsensus module was used to compare each Genewise model against mouse cDNA and EST alignments in the region (see next section on how these alignments were generated), where exons in the Genewise model were scored for overlapping with exons of cDNA/EST alignments, and model(s) with the highest combined score in a region were kept. The generation of transcript models using data from related species is referred to as the "Similarity stage" [Figure 3]. This stage resulted in 7525 and 1267 coding models supported by mammalian Uniprot proteins and non-mammalian vertebrate Uniprot proteins respectively.


Figure 3: Alignment and filtering of mammalian and vertebrate proteins.

cDNA and EST Alignment

Approximate time: 1-2 weeks

Mouse cDNAs and ESTs were downloaded from ENA/Genbank/DDBJ and RefSeq [9.], clipped to remove polyA tails, and aligned to the genome using Exonerate (*est2genome* model) [Figure 4].
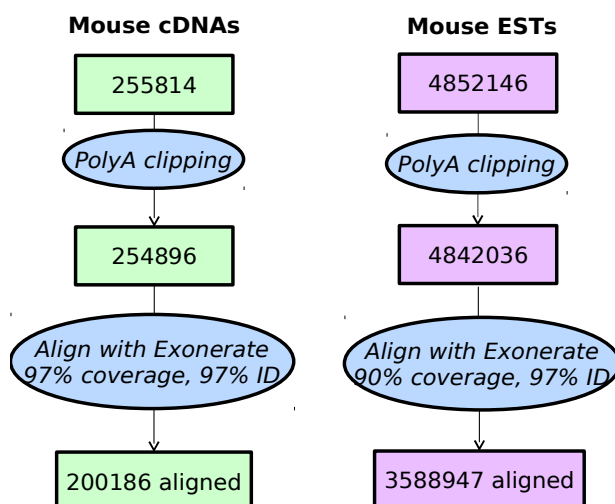
Mouse cDNAs   **Mouse ESTs**

```
  255814                  4852146

PolyA clipping          PolyA clipping

  254896                  4842036

Align with Exonerate    Align with Exonerate
97% coverage, 97% ID    90% coverage, 97% ID

200186 aligned          3588947 aligned
```

Figure 4: Alignment of mouse cDNAs and ESTs to the mouse genome.

200186 (of 255814) mouse cDNAs aligned and 358847 (of 4852146) mouse ESTs aligned. All alignments were at a cut-off of 97% identity. The coverage cut-off for cDNA alignments (97%) was higher than that for ESTs (90%) because cDNAs are generally less fragmented than ESTs. EST alignments were used to generate EST-based gene models similar to those for human [14.] and these are displayed on the website in a separate track from the Ensembl gene set.

Filtering Coding Models

Approximate time: 2 weeks

The set of coding models was finalised after another stage of filtering, which involved manual removal of some more Targetted models supported by dubious mouse protein/cDNA evidence on a case-by-case basis, and removal of ~60% of Similarity alignments which contained non-canonical (non GT/AG) splice sites using a Perl script. The Apollo software [15.] was used to visualise the results of filtering.

Addition of UTR to coding models

Approximate time: 2 weeks

After finalising the set of coding models, those generated by Genewise alignments were extended into the untranslated regions (UTRs) using mouse cDNAs. (Coding models generated by Exonerate's *cdna2genome* model already contained UTR annotations and hence did not go through this UTR addition step.) Where available, mouse ditag alignments were used to guide the positioning of UTRs and add additional weight to some UTR structures, while RefSeq "NM" cDNA vs "NP" protein pairing information was used to ensure the correct matching of cDNAs to coding models supported by RefSeq proteins. This resulted in 28266 (of 36726) coding models from 25984 mouse proteins with UTR, and 272 (of 3757) coding models from 114 Uniprot proteins with UTR.

Generating multi-transcript Ensembl genes

Approximate time: 3-4 weeks

The above steps generated a large set of potential transcript models (with or without UTR), many of which overlapped one another. Redundant transcript models were collapsed and the remaining unique set of transcript models were clustered into multi-transcript genes where each transcript in a gene has at least one coding exon that overlaps a coding exon from another transcript within the same gene. The resulting Ensembl gene set contained 23987 genes, of which 23113 contained transcripts supported by mouse cDNAs/proteins only (from the "Targetted" stage of the build), and 874 contained transcripts supported by Uniprot proteins only (from the "Similarity" stage of the build). [Figure 5].   The Ensembl genes were associated with a total of 37361 Ensembl transcripts, of which 36484 were supported by mouse cDNAs/proteins, and 877 had support from Uniprot proteins [Figure 6].

**Figure 5: Supporting evidence for mouse Ensembl gene set.**

Figure 6: Supporting evidence for mouse Ensembl transcript set.

Pseudogenes, immunoglobulin genes, mitochondrial genes

Approximate time: 3 weeks

The Ensembl gene set was screened for pseudogenes and retrotransposed genes. Next, mouse immunoglobulin (Ig) genes were annotated using the Ensembl "Ig genebuild" pipeline [16.]. Briefly, mouse proteins and cDNAs for Ig genes were downloaded from IMGT [17.] and aligned to the mouse genome using Exonerate. The Exonerate alignments were processed to join the V/D/J/C segments together into Ig gene models, which were then compared to the Ig genes already present in the Ensembl gene set (generated at the Targetted stage).   If the models generated by the "Ig genebuild" pipeline overlapped with existing Ensembl genes at the exon level,

the existing Ensembl genes will be replaced by the new Ig gene models, for the latter are usually more accurate representations of Ig genes.   Also imported into the Ensembl gene set were annotation of mitochondrial genes in INDSC [18.]   and short non-coding   RNAs (e.g. miRNAs, snoRNAs) generated by the ncRNA pipeline [19.].

Merging Ensembl and Vega gene sets, annotating long intergenic non-coding RNA genes and generating the final gene set.

Approximate time: 10 weeks

Following the completion of the Ensembl gene set, Ensembl annotations and manual annotations (primarily generated by the HAVANA team at the Wellcome Trust Sanger Institute) from the Vega database [20.,21.] were merged at the transcript level to create the final gene set.   The Vega database (as of 17 September 2010) contained 20773 genes and 65720 transcripts.   In the merge process, Ensembl and Vega transcripts were merged if they had identical exon-intron structures. If transcripts from the two annotation sources matched at all internal exon-intron boundaries, i.e. had identical splicing pattern, but one of them had longer terminal exon(s) (usually the UTRs), they were merged too, but the resulting merged transcript would adopt the exon-intron structure of the Vega transcript, for we prioritised Vega annotation over Ensembl's. Transcripts which had not been merged, either because of differences in internal exon-intron boundaries or presence of transcripts in only one annotation source, were transferred from the source to the final gene set intact.

The Ensembl-Vega merge code also took into account the biotype and supporting evidence associated with the transcripts from both annoation sources. For a pair of transcripts to be merged, if there was a mismatch in biotype, e.g. the Ensembl transcript is protein-coding but the Vega counterpart is non-coding, the Vega biotype would have precedence over Ensembl's, and

the Ensembl transcript would undergo a biotype change to match its Vega counterpart. Ensembl transcript's translation would be removed too if the transcript has lost its protein-coding biotype. Biotype conflicts between Ensembl and Vega were always reported to the HAVANA team for investigation, and when resolved, could improve the merged gene set in the future. As for supporting evidence, the merge of Ensembl and Vega transcripts also involved merging of protein/cDNA supporting evidence associated with the transcripts to ensure the basis on which the annotations were made would not be lost.

Following the merge, long intergenic non-coding RNA genes (lincRNAs) were annotated by the Ensembl lincRNA pipeline [19.] and incorporated in the final gene set.

An important feature of the merged gene set is the presence of all Vega source transcripts. This has been made possible by allowing Vega annotation to take precedence over Ensembl's when merging transcripts which do not match at their terminal exons or have different biotypes. Of all Vega source transcripts, about 27% of them merged with Ensembl transcripts. The vast majority of merged transcripts (91.2%) are of protein-coding biotype. Vega transcripts which were not merged (~73% of Vega source transcripts) were mostly alternative splice variants and/or non-coding, and were carried over into the final gene set intact. The final gene set consists of 36817 genes and 93809 transcripts. Of the 93809 transcripts, 18.96% (17788) were the result of merging Ensembl and Vega annotations, 22.18% (20809) originated from Ensembl, 52.04% (48821) originated from Vega, and a remaining ~6.8% were incorporated from other sources (e.g. immunoglobulin gene segments/transcripts imported from IMGT data).

As a quality-control measure, Ensembl translations of protein-coding transcripts in the final merged gene set were aligned against the NCBI RefSeq and Uniprot/SwissProt sets of public curated protein sequences (which were used in the "Targetted" stage of the gene build) to calculate the proportion of curated sequences covered by the merged gene set. O99% of

RefSeq and SwissProt proteins were represented in the merged gene set, and in the majority of cases, there was a 100% match between the curated protein and Ensembl translation.

Protein annotation, Cross-referencing, Stable Identifiers

Approximate time: 4 weeks

Before public release the transcripts and translations were given external references (cross-references to external databases), while translations were searched for domains/signatures of interest and labelled where appropriate. Stable identifiers were assigned to each gene, transcript, exon and translation. (When annotating a species for the first time, these identifiers are auto-generated. In all subsequent annotations for a species, the stable identifiers are propagated based on comparison of the new gene set to the previous gene set.)

Further information on the Ensembl gene set

The main focus of the Ensembl automatic gene annotation pipeline is to generate a conservative set of protein-coding gene models, although some non-coding genes and pseudogenes may also annotated. The Vega project [20., 21.], on the other hand, focuses on manually annotating alternative splice variants for all genes and annotating a much wider range of gene/transcript types, including non-coding genes (e.g. processed transcripts, nonsense-mediated decay transcripts, polymorphic pseudogenes) [22.] Therefore, the Ensembl and Vega annotation approaches complement each other and by merging the Ensembl and Vega annotations, we aim to provide a more comprehensive final gene set for mouse.

Every gene model produced by the Ensembl gene annotation pipeline is supported by biological sequence evidence (see the "Supporting evidence" link on the left-hand menu of a Gene page or Transcript page); *ab initio* models are not included in our gene set. *Ab initio* predictions and the full set of cDNA and EST alignments to the genome are available on our website.

The quality of a gene set is dependent on the quality of the genome assembly. Genome assembly can be assessed in a number of ways, including:

1. Coverage estimate
   - A higher coverage usually indicates a more complete assembly.
   - Using Sanger sequencing only, a coverage of at least 2x is preferred.

2. N50 of contigs and scaffolds
   - A longer N50 usually indicates a more complete genome assembly.
   - Bearing in mind that an average human gene may be 10-15 kb in length, contigs shorter than this length will be unlikely to hold full-length gene models.

3. Number of contigs and scaffolds
   - A lower number toplevel sequences usually indicates a more complete genome assembly.

4. Alignment of cDNAs and ESTs to the genome
   - A higher number of alignments, using stringent thresholds, usually indicates a more complete genome assembly.

More information on the Ensembl automatic gene annotation process can be found at:

- Curwen V, Eyras E, Andrews TD, Clarke L, Mongin E, Searle SM, Clamp M. **The Ensembl automatic gene annotation system.** *Genome Res.* 2004, **14(5):**942-50. [PMID: 15123590]
- Potter SC, Clarke L, Curwen V, Keenan S, Mongin E, Searle SM, Stabenau A, Storey R, Clamp M. **The Ensembl analysis pipeline.** *Genome Res.* 2004, **14(5):**934-41. [PMID: 15123589]
- http://www.ensembl.org/info/docs/genebuild/genome_annotation.html
- http://cvs.sanger.ac.uk/cgi-bin/viewvc.cgi/ensembl-doc/pipeline_docs/the_genebuild_process.txt?root=ensembl&view=co

References

1. Smit, AFA, Hubley, R & Green, P: **RepeatMasker Open-3.0.** 1996-2010. www.repeatmasker.org

2. Kuzio J, Tatusov R, and Lipman DJ: **Dust.** Unpublished but briefly described in: Morgulis A, Gertz EM, Schäffer AA, Agarwala R. A Fast and Symmetric DUST Implementation to Mask Low-Complexity DNA Sequences. *Journal of Computational Biology* 2006, **13(5):**1028-1040.

3. Benson G. **Tandem repeats finder: a program to analyze DNA sequences.** *Nucleic Acids Res.* 1999, **27(2):**573-580. [PMID: 9862982]. http://tandem.bu.edu/trf/trf.html

4. Down TA, Hubbard TJ: **Computational detection and location of transcription start sites in mammalian genomic DNA.** *Genome Res.* 2002 **12(3):**458-461. http://www.sanger.ac.uk/resources/software/eponine/ [PMID: 11875034]

5. Davuluri RV, Grosse I, Zhang MQ: **Computational identification of promoters and first exons in the human genome.** *Nat Genet.* 2001, **29(4):**412-417. [PMID: 11726928]

6. Lowe TM, Eddy SR: **tRNAscan-SE: a program for improved detection of transfer RNA genes in genomic sequence.** *Nucleic Acids Res.* 1997, **25(5):**955-64. [PMID: 9023104]

7. Burge C, Karlin S: **Prediction of complete gene structures in human genomic DNA.** *J Mol Biol.* 1997, **268(1):**78-94. [PMID: 9149143]

8. Goujon M, McWilliam H, Li W, Valentin F, Squizzato S, Paern J, Lopez R: **A new bioinformatics analysis tools framework at EMBL-EBI. Nucleic Acids Res.** 2010, **38 Suppl:**W695-699. http://www.uniprot.org/downloads [PMID: 20439314]

9. Sayers EW, Barrett T, Benson DA, Bolton E, Bryant SH, Canese K, Chetvernin V, Church DM, Dicuccio M, Federhen S, Feolo M, Geer LY, Helmberg W, Kapustin Y, Landsman D, Lipman DJ, Lu Z, Madden TL, Madej T, Maglott DR, Marchler-Bauer A, Miller V, Mizrachi I, Ostell J, Panchenko A, Pruitt KD, Schuler GD, Sequeira E, Sherry ST, Shumway M, Sirotkin K, Slotta D, Souvorov A, Starchenko G, Tatusova TA, Wagner L, Wang Y, John Wilbur W, Yaschenko E, Ye J: **Database resources of the National Center for Biotechnology Information.** *Nucleic Acids Res.* 2010, **38(Database issue):D5-16.** [PMID: 19910364]

10. http://www.ebi.ac.uk/ena/

11. Altschul SF, Gish W, Miller W, Myers EW, Lipman DJ: **Basic local alignment search tool.** *J Mol Biol.* 1990, **215(3):**403-410. [PMID: 2231712.]

12. Slater GS, Birney E: **Automated generation of heuristics for biological sequence comparison.** *BMC Bioinformatic*s 2005, **6:**31. [PMID: 15713233]

13.   Birney E, Clamp M, Durbin R: **GeneWise and Genomewise.** *Genome Res.* 2004, **14(5):**988-995. [PMID: 15123596]

14.   Eyras E, Caccamo M, Curwen V, Clamp M. **ESTGenes: alternative splicing from ESTs in Ensembl.** *Genome Res.* 2004 **14(5):**976-987. [PMID: 15123595]

15.   Lewis SE, Searle SM, Harris N, Gibson M, Lyer V, Richter J, Wiel C, Bayraktaroglir L, Birney E, Crosby MA, Kaminker JS, Matthews BB, Prochnik SE, Smithy CD, Tupy JL, Rubin GM, Misra S, Mungall CJ, Clamp ME: **Apollo: a sequence annotation editor.** *Genome Biol.* 2002, **3(12):**RESEARCH0082. [PMID: 12537571]

16.   http://www.ensembl.org/info/docs/genebuild/ig_tcr.html

17.   ftp://ftp.cines.fr/IMGT/IMGT.zip

18.   http://www.ncbi.nlm.nih.gov/nuccore/NC_005089

19.   http://www.ensembl.org/info/docs/genebuild/ncrna.html

20.   http://vega.sanger.ac.uk/Mus_musculus/Info/Index

21.   L. G. Wilming, J. G. R. Gilbert, K. Howe, S. Trevanion,T. Hubbard and J. L. Harrow: The vertebrate genome annotation (Vega) database. Nucleic Acid Res. 2008 Jan; Advance Access published on November 14, 2007; doi:10.1093/nar/gkm987

22.   http://vega.sanger.ac.uk/info/about/gene_and_transcript_types.html